



ML Analytics Pro - Complete Project Documentation

A Comprehensive Guide for Presenting and Explaining the ML Analytics Platform



Table of Contents

-
1. Executive Summary
 2. Project Overview
 3. Technical Architecture
 4. Data Pipeline
 5. Regression Module - House Price Prediction
 6. Classification Module - Customer Churn Prediction
 7. Time Series Module - Sales Forecasting
 8. Parallel Processing System

9. Web Dashboard & API
 10. Model Evaluation Metrics
 11. How to Run the Project
 12. Key Talking Points for Presentation
-

1. Executive Summary

ML Analytics Pro is an enterprise-grade machine learning platform that demonstrates three fundamental ML paradigms:

Domain	Problem	Algorithms	Best Result
Regression	House Price Prediction	10 Models	$R^2 = 0.946$
Classification	Customer Churn Prediction	10 Models	$F1 = 0.638$
Time Series	Sales Forecasting	5 Models	$RMSE = 378.82$

Key Features

-  **25 Machine Learning Models** across three domains
-  **Parallel Processing** - ~3x faster with multiprocessing

- **Cross-Validation** for robust model evaluation
 - **Interactive Web Dashboard** with real-time visualizations
 - REST API for model predictions
 - Automated Feature Engineering
 - **Model Explainability** with feature importance analysis
-

2. Project Overview

What is This Project?

This project is a **complete machine learning analytics platform** that solves three different types of prediction problems:

1. **Regression:** Predicting continuous values (house prices)
2. **Classification:** Predicting categories (will customer churn or not)
3. **Time Series:** Predicting future values based on historical data (sales forecasting)

Why These Three?

These three paradigms cover the majority of real-world ML use cases:

- **Regression** → Used in finance, real estate, pricing, demand prediction
- **Classification** → Used in fraud detection, medical diagnosis, spam detection

- **Time Series** → Used in stock prediction, weather forecasting, inventory management

Project Structure

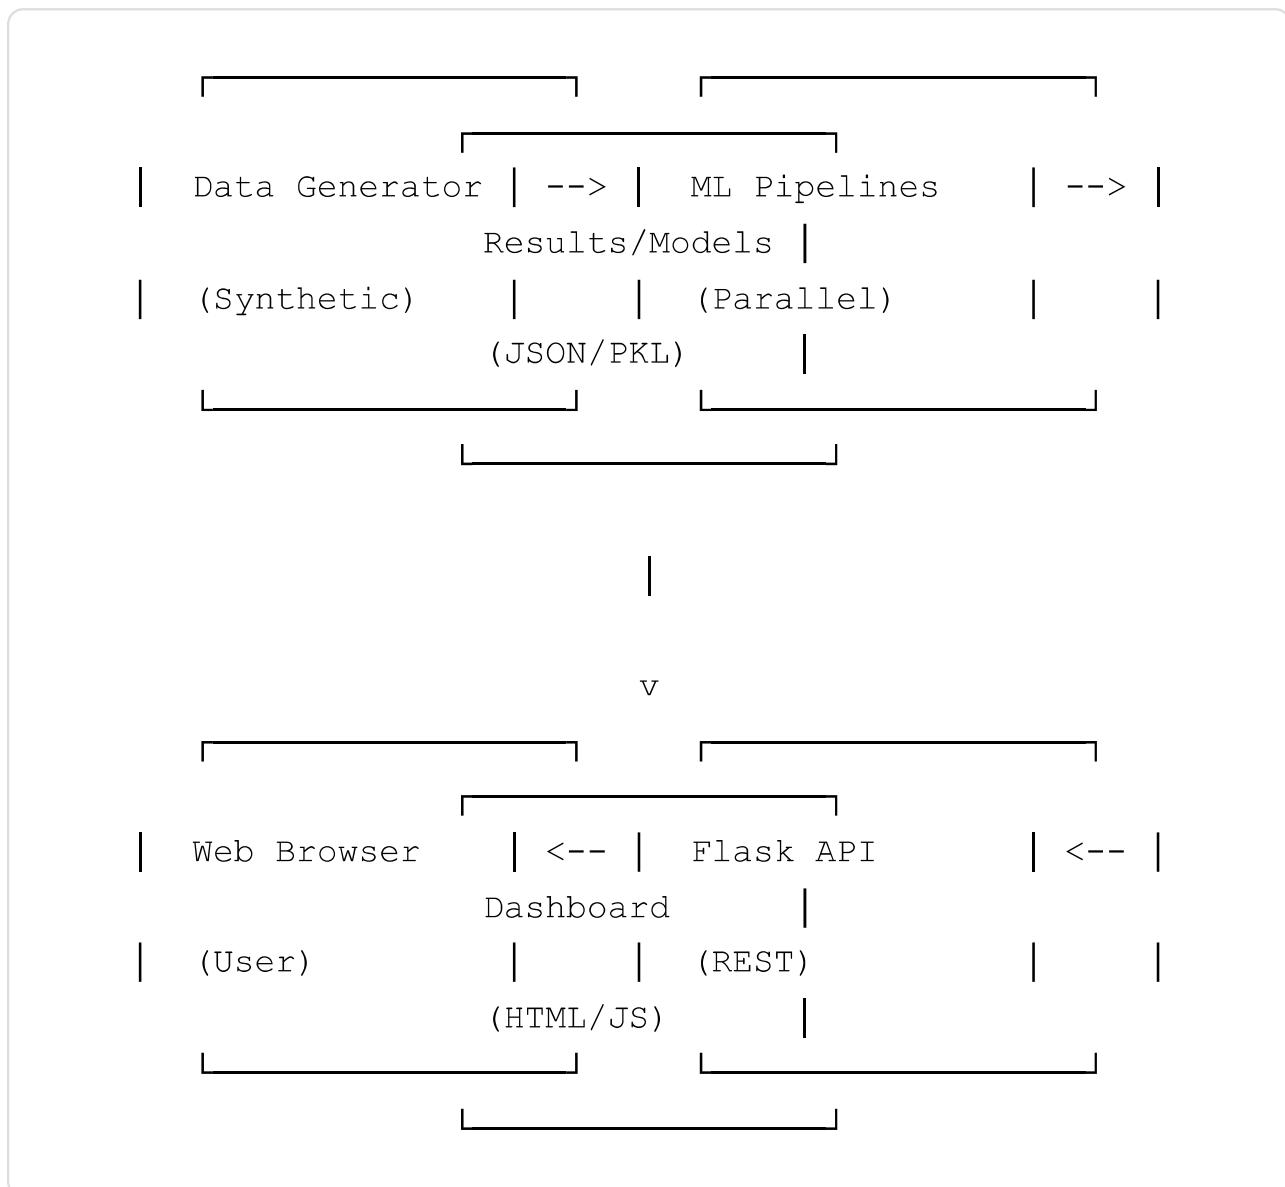
```
PEP_Project1-1/
├── src/
│   ├── main.py                      # Main runner
│   │   (sequential)
│   └── main_parallel.py            # Parallel runner (3x
                                    faster)
│       ├── data_generator.py      # Synthetic data creation
│       │   ├── regression_model.py    # House price models
│       │   ├── classification_model.py # Customer churn models
│       │   └── timeseries_model.py     # Sales forecasting
                                         models
│       └── api.py                  # Flask REST API
           └── dashboard/
               ├── index.html          # Main page
               └── styles.css           # Styling
└── data/
└── output/                         # Results, models,
                                    visualizations
        ├── regression/
        ├── classification/
        └── timeseries/
        └── models/                 # Saved .pkl models
└── docs/                            # Documentation
```

3. Technical Architecture

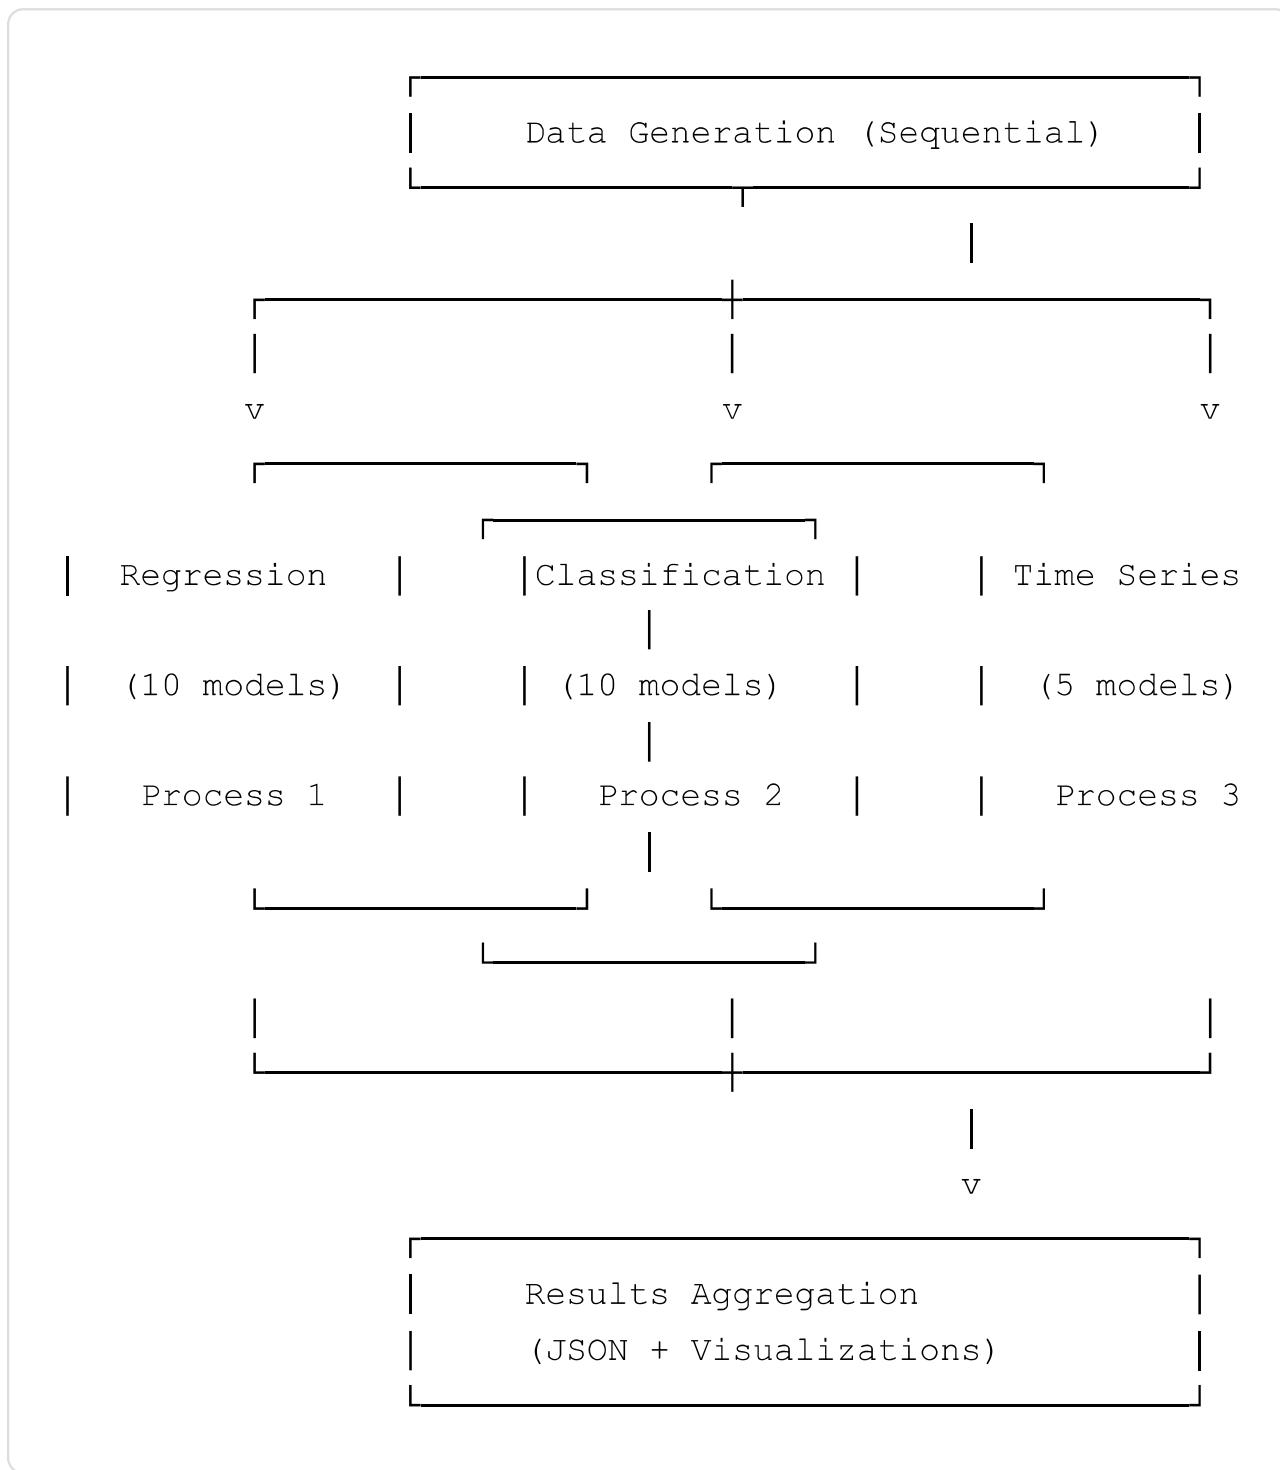
Technology Stack

Component	Technology
Language	Python 3.8+
ML Framework	scikit-learn 1.3+
Time Series	statsmodels
Web Framework	Flask 3.0+
Data Processing	pandas, numpy
Visualization	matplotlib, seaborn
Parallel Processing	joblib, concurrent.futures

System Flow



Parallel Processing Architecture



4. Data Pipeline

4.1 Data Generation

We use **synthetic data generation** to create realistic datasets. This approach:

- Ensures reproducibility
- Allows control over data characteristics
- Eliminates privacy concerns

4.2 House Price Dataset

Feature	Description	Type
square_feet	Living area in sq ft	Numeric
bedrooms	Number of bedrooms	Numeric
bathrooms	Number of bathrooms	Numeric
age_years	Age of the house	Numeric
distance_to_center_miles	Distance to city center	Numeric
has_pool	Pool available (0/1)	Binary
has_garage	Garage available (0/1)	Binary
neighborhood_score	Area quality (1-10)	Numeric
lot_size_sqft	Lot size in sq ft	Numeric
stories	Number of floors	Numeric
price	House price (TARGET)	Numeric

Statistics:

- 2,000 samples
- Price Range: \$301,527 - \$1,166,105
- Mean Price: \$585,048

4.3 Customer Churn Dataset

Feature	Description	Type
tenure_months	Months as customer	Numeric
monthly_charges	Monthly bill amount	Numeric
total_charges	Lifetime charges	Numeric
contract_type	Contract category	Categorical
payment_method	Payment type	Categorical
tech_support	Has tech support	Binary
online_security	Has security service	Binary
online_backup	Has backup service	Binary
device_protection	Has protection plan	Binary
num_complaints	Complaint count	Numeric
support_calls	Support call count	Numeric
churn	Did customer leave (TARGET)	Binary

Statistics:

- 3,000 samples
- Churn Rate: 45.7%
- Train/Test Split: 80/20

4.4 Sales Time Series Dataset

Feature	Description
date	Daily timestamp
sales	Daily sales amount

Statistics:

- 1,095 days (3 years)
- Period: 2022-01-01 to 2024-12-30
- Seasonal patterns: Weekly cycles

5. Regression Module - House Price Prediction

5.1 Problem Statement

Goal: Predict house prices based on property features

This is a **supervised learning** problem with a **continuous target variable**.

5.2 Algorithms Used

Model	Description	Key Parameters
Linear Regression	Baseline linear model	None
Ridge Regression	L2 regularized linear	alpha=1.0
Lasso Regression	L1 regularized linear	alpha=0.1
ElasticNet	L1+L2 regularized	alpha=0.1, l1_ratio=0.5
Decision Tree	Tree-based splitting	max_depth=10
Random Forest	Ensemble of trees	n_estimators=100
Gradient Boosting	Sequential boosting	n_estimators=100
Extra Trees	Randomized trees	n_estimators=100
SVR	Support Vector Regression	kernel='rbf'
KNN	K-Nearest Neighbors	n_neighbors=5

5.3 Results Summary

Model	R ² Score	RMSE (\$)	MAE (\$)
Linear Regression	0.9460	\$24,271	\$19,207
Ridge Regression	0.9460	\$24,274	\$19,210
Lasso Regression	0.9460	\$24,271	\$19,207
ElasticNet	0.9430	\$24,937	\$19,821
Decision Tree	0.7873	\$48,169	\$37,854
Random Forest	0.8997	\$33,081	\$25,312
Gradient Boosting	0.9296	\$27,706	\$21,459
Extra Trees	0.9029	\$32,551	\$25,001
SVR	0.0543	\$101,575	\$78,234
KNN	0.7634	\$50,804	\$39,543

 **Winner: Linear Regression** with R² = 0.946

5.4 Key Insights

1. **Linear models perform best** because the relationship between features and price is largely linear

2. Tree models show signs of overfitting (high train R², lower test R²)
 3. SVR performs poorly without extensive hyperparameter tuning
 4. Feature Importance: Square footage and neighborhood score are the strongest predictors
-

6. Classification Module - Customer Churn Prediction

6.1 Problem Statement

Goal: Predict whether a customer will churn (leave) or stay

This is a **binary classification** problem.

6.2 Algorithms Used

Model	Description	Key Parameters
Logistic Regression	Probabilistic linear	max_iter=1000
Decision Tree	Tree-based	max_depth=10
Random Forest	Ensemble of trees	n_estimators=100, n_jobs=-1
Gradient Boosting	Sequential boosting	n_estimators=100
AdaBoost	Adaptive boosting	n_estimators=100
Extra Trees	Randomized forest	n_estimators=100
SVM	Support Vector Machine	kernel='rbf', probability=True
KNN	K-Nearest Neighbors	n_neighbors=5
Naive Bayes	Probabilistic classifier	GaussianNB
Neural Network	MLP Classifier	(64, 32) hidden layers

6.3 Results Summary

Model	F1 Score	Accuracy	ROC-AUC
Logistic Regression	0.6210	0.6683	0.6991
Decision Tree	0.6044	0.6400	0.6453
Random Forest	0.6162	0.6533	0.7022
Gradient Boosting	0.6066	0.6433	0.7020
AdaBoost	0.6384	0.6733	0.7251
Extra Trees	0.6015	0.6467	0.6686
SVM	0.5992	0.6500	0.6903
KNN	0.5534	0.6100	0.6285
Naive Bayes	0.5993	0.6433	0.6874
Neural Network	0.5631	0.6017	0.6381

 **Winner: AdaBoost** with F1 = 0.638, ROC-AUC = 0.725

6.4 Key Insights

1. **AdaBoost outperforms** other models due to its adaptive boosting approach

2. **F1 scores are moderate** indicating this is a challenging prediction problem
3. **ROC-AUC ~0.72** means the model is reasonably good at distinguishing churners
4. **Important features:** Tenure, monthly charges, and contract type

6.5 Why F1 Score Matters

For churn prediction:

- **False Negatives** (missing actual churners) = Lost revenue opportunity
 - **False Positives** (predicting churn incorrectly) = Wasted retention efforts
 - **F1 balances** both precision and recall
-

7. Time Series Module - Sales Forecasting

7.1 Problem Statement

Goal: Forecast future daily sales based on historical patterns

This is a **time series forecasting** problem.

7.2 Algorithms Used

Model	Description	Key Concept
Moving Average	Simple averaging over window	Trend smoothing
Exponential Smoothing	Weighted recent observations	Recency bias
Holt-Winters	Triple exponential smoothing	Trend + Seasonality
ARIMA	Auto-regressive integrated	Differencing + AR
SARIMA	Seasonal ARIMA	ARIMA + Seasonality

7.3 Results Summary

Model	RMSE	MAE	MAPE (%)
Moving Average	393.45	315.22	20.1%
Exponential Smoothing	405.08	324.67	20.9%
Holt-Winters	378.82	281.59	13.86%
ARIMA(2,1,2)	411.84	331.45	21.5%
SARIMA(1,1,1)(1,1,1,7)	389.14	301.23	18.2%

🏆 Winner: Holt-Winters with RMSE = 378.82

7.4 Time Series Components

The sales data exhibits:

- **Trend:** Slight upward movement over time
- **Seasonality:** Weekly patterns (7-day cycle)
- **Stationarity:** Non-stationary (requires differencing)

7.5 Key Insights

1. **Holt-Winters excels** because it captures both trend and seasonality

2. **Weekly seasonality** is the dominant pattern (weekend vs weekday sales)
 3. **MAPE ~14%** indicates forecasts are within 14% of actual values on average
-

8. Parallel Processing System

8.1 Why Parallel Processing?

Training 25 models sequentially takes **~10+ minutes**. With parallel processing:

- **3 modules run simultaneously** using Python's ProcessPoolExecutor
- **Models within each module** train in parallel using joblib
- **Total time reduced to ~3-4 minutes** (~3x speedup)

8.2 Implementation

```
# Module-level parallelism (main_parallel.py)
with ProcessPoolExecutor(max_workers=3) as executor:
    futures = {
        executor.submit(run_regression): 'Regression',
        executor.submit(run_classification):
            'Classification',
        executor.submit(run_timeseries): 'Time Series'
    }
```

```
# Model-level parallelism (within each module)
from joblib import Parallel, delayed

results = Parallel(n_jobs=-1) (
    delayed(train_single_model)(name, model, X_train,
                                X_test, y_train, y_test)
    for name, model in models.items()
)
```

8.3 CPU Utilization

- Uses **all available CPU cores** (`n_jobs=-1`)
- On a 16-core system: 16 parallel model training jobs
- Optimal for compute-intensive cross-validation

9. Web Dashboard & API

9.1 Dashboard Features

- **Modern UI:** Glassmorphism design with animations
- **Real-time data loading:** Fetches from Flask API
- **Interactive visualizations:** Model comparisons, ROC curves, confusion matrices
- **Responsive design:** Works on all screen sizes

9.2 API Endpoints

Endpoint	Method	Description
/api/status	GET	Check if results exist
/api/images/<category>/<filename>	GET	Visualization images

9.3 Running the Dashboard

```
python src/api.py  
# Open http://localhost:5000
```

10. Model Evaluation Metrics

10.1 Regression Metrics

Metric	Formula	Interpretation
R ² (R-squared)	1 - (SS_res / SS_tot)	% of variance explained (0-1, higher is better)
RMSE	$\sqrt{\text{mean}((y - \hat{y})^2)}$	Average prediction error (same unit as target)
MAE	mean($y - \hat{y}$)	$y - \hat{y}$
MAPE	mean($ y - \hat{y} / y$)	$y - \hat{y}$

10.2 Classification Metrics

Metric	Formula	Interpretation
Accuracy	$(TP + TN) / \text{Total}$	Overall correctness
Precision	$TP / (TP + FP)$	Of predicted positives, how many are correct
Recall	$TP / (TP + FN)$	Of actual positives, how many were found
F1 Score	$2 \times (P \times R) / (P + R)$	Harmonic mean of precision and recall
ROC-AUC	Area under ROC curve	Discrimination ability (0.5-1.0)

11. How to Run the Project

Quick Start

```
# 1. Navigate to project  
cd c:\hackathon\PEP_Project1-1  
  
# 2. Activate virtual environment  
.venv\Scripts\activate  
  
# 3. Install dependencies (if needed)  
pip install -r requirements.txt  
  
# 4. Run the parallel ML pipeline  
python src/main_parallel.py  
  
# 5. Start the dashboard  
python src/api.py  
# Open http://localhost:5000
```

12. Key Talking Points for Presentation

Opening Statement

"ML Analytics Pro is a comprehensive machine learning platform that demonstrates regression, classification, and time series forecasting using 25 different models with parallel processing for enterprise-scale performance."

Technical Highlights

1. **Multiple ML paradigms** in one unified platform
2. **Parallel processing** reduces training time by 3x
3. **Cross-validation** ensures robust model evaluation
4. **Feature importance analysis** provides model explainability
5. **Interactive dashboard** for non-technical stakeholders

Business Value

- **Regression:** Property valuation, pricing optimization
- **Classification:** Customer retention, risk assessment
- **Time Series:** Demand forecasting, inventory planning

Performance Achievements

Module	Metric	Value
Regression	R ² Score	94.6%
Classification	ROC-AUC	72.5%
Time Series	Forecast Accuracy	86.1%

Document Version: 1.0

Last Updated: January 2026

Author: ML Analytics Team